

Homework 3

Due by **12:30pm, Friday, September 30, 2016.**

Make sure you follow all the homework policies (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/policies/hw-policy.html>).

All submissions should be done via Autolab (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/autolab.html>).

The support page for matrix vector multiplication (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/support/matrix-vect/index.html>) should be very useful for this homework.

1. This might seem counter-intuitive but actually in this case looking at specific examples of these structured matrices might distract you from the high level structure in these matrices, which is more apparent in the formal definition. The final algorithm is easiest to design if you just concentrate on this abstract high level structure.

Sample Problem

The Problem

For this and the remaining problems, we will be working with $n \times n$ matrices (or two-dimensional arrays). So for example the following is a 3×3 matrix

$$\mathbf{M} = \begin{pmatrix} 1 & 2 & -3 \\ 2 & 9 & 0 \\ 6 & -1 & -2 \end{pmatrix}.$$

Given an $n \times n$ matrix \mathbf{A} , we will refer to the entry in the i th row and j th column (for $0 \leq i, j < n$) as $\mathbf{A}[i][j]$. So for example, for the matrix \mathbf{M} above, $\mathbf{M}[1][2] = 0$. (We will use the convention that the top left element is the $[0][0]$ element.) We will also work with vectors of length n , which are just arrays of size n . For example the following is a vector of length 3

$$\mathbf{z} = \begin{pmatrix} 2 \\ 3 \\ -1 \end{pmatrix}.$$

As usual we will use $x[i]$ to denote the i th entry in the array/vector. For example, in the vector/array \mathbf{z} above, we have $\mathbf{z}[2] = -1$.

We are finally ready to define the matrix-vector multiplication problem. Given an $n \times n$ matrix \mathbf{A} and a vector \mathbf{x} of length n , their product is denoted by

$$\mathbf{y} = \mathbf{A} \cdot \mathbf{x},$$

where \mathbf{y} is also a vector of length n and its i th entry for $0 \leq i < n$ is defined as follows:

$$y[i] = \sum_{j=0}^{n-1} \mathbf{A}[i][j] \cdot \mathbf{x}[j].$$

For example, here is the worked out example for \mathbf{M} and \mathbf{z} above:

$$\begin{pmatrix} 1 & 2 & -3 \\ 2 & 9 & 0 \\ 6 & -1 & -2 \end{pmatrix} \cdot \begin{pmatrix} 2 \\ 3 \\ -1 \end{pmatrix} = \begin{pmatrix} 1 \times 2 + 2 \times 3 + (-3) \times (-1) \\ 2 \times 2 + 9 \times 3 + 0 \times (-1) \\ 6 \times 2 + (-1) \times 3 + (-2) \times (-1) \end{pmatrix} = \begin{pmatrix} 11 \\ 31 \\ 11 \end{pmatrix}.$$

Finally, we come to the actual question:

1. Design an $O(n^2)$ algorithm, which given any $n \times n$ matrix \mathbf{A} and vector \mathbf{x} of length n correctly computes $\mathbf{A} \cdot \mathbf{x}$.

Note

This is exactly Exercise 1 in the support page on matrix vector multiplication ([../support/matrix-vect/index.html](http://www-student.cse.buffalo.edu/~atri/cse331/fall16/support/matrix-vect/index.html)).

Hint

The obvious algorithm works for this problem.

2. Argue that your algorithm in part above runs in time $\Omega(n^2)$. (Hence, conclude that your algorithm runs in time $\Theta(n^2)$.)
3. (**Only think about this if you have time to spare.**) Argue that *any* algorithm that solves the matrix vector multiplication problem (for arbitrary \mathbf{A} and \mathbf{x}) needs to take $\Omega(n^2)$ time.

Note

This is the same as Exercise 2 in the support page on matrix vector multiplication ([../support/matrix-vect/index.html](http://www-student.cse.buffalo.edu/~atri/cse331/fall16/support/matrix-vect/index.html)). Further, note that this part implies the second part but *if* you want to use third part to solve the second part, then you will have to present the solution to third part too. I would recommend that you solve part three only if you want to stretch your algorithmic thinking.

[Click here for the Solution](#)

Submission

You will **NOT** submit this question. This is for you to get into thinking more about matrix vector multiplication.

Question 1 (Programming Assignment) [40 points]

</> Note

This assignment can be solved in either Java, Python or C++ (you should pick the language you are most comfortable with). Please make sure to look at the supporting documentation and files for the language of your choosing.

The Problem

In this problem we will consider matrix and vector multiplication, but with the caveat that we already know something about the matrix. In particular for any $0 \leq i, j \leq n - 1$, we have

$$\mathbf{U}_n[i][j] = \begin{cases} 0 & \text{if } j < i \\ 1 & \text{otherwise.} \end{cases}$$

If you "draw" this matrix you will note that all of the elements below the diagonal is zero: such matrices are called *upper triangular matrices*. Further, all non-zero values (in the "upper triangular" part) are all 1. Note that we only need to know the value of n to completely determine \mathbf{U}_n .

You have to write code to solve the following computational problem: given any integer $n \geq 1$ and a vector \mathbf{x} of length n , return the vector $\mathbf{y} = \mathbf{U}_n \cdot \mathbf{x}$. Read on for more details on the format etc.

Sometimes it is easier to visualize a specific \mathbf{U}_n . Below is \mathbf{U}_5 written down in a 2D-array form:

1	1	1	1	1
0	1	1	1	1
0	0	1	1	1
0	0	0	1	1
0	0	0	0	1

Input

You will **not** receive the matrix as the input.

</> Note

You *could* build the matrix \mathbf{U}_n yourself but this is **HIGHLY DISCOURAGED**. For the biggest input, you will need about 1 TB to store this matrix explicitly!

What you will receive is the vector itself. Each input file will contain one vector. The length of the vector will be on the first line of the file. The vector itself will take up the remaining lines. Once again, you will not have to handle the file reading and the creation of data structures. The packaged driver will handle that for you.

The input file can be summed up in the following format:

```
n
x0
x1
x2
.
.
.
xn-1
```

where each x_k is a value of the vector, in vertical order.

This is an example of what a file may look like, for size five:

```
5
43
234
-12
32
-44
```

Output

The output will be the vector \mathbf{y} in vertical order. I.e.

```
y0
y1
y2
.
.
.
yn-1
```

The correct result for the earlier example of \mathbf{x} would be:

```
253
210
-24
-12
-44
```

Hint

The best possible algorithm for this problem runs in time $O(n)$.

! Note

Both the input and output parsers in each of the three languages are already written for you.

Note that you have to work with the input data structures provided (which will come pre-loaded with the data from an input file).

! Addition is the only change you should make

Irrespective of what language you use, you will have to submit just one file. That file will come pre-populated with some stuff in it. You should **not** change any of those things because if you do you might break what the grader expects and end up with a zero on the question. You should of course add stuff to it (including helper functions and data structures as you see fit).

Java

Python

C++

[Download Java Skeleton Code \(HW3Java.zip\)](#)

Directory Structure

```

├── Driver.java
├── HW3_Student_Solution.java
├── testcases/
│   ├── input1.txt
│   ├── input2.txt
│   ├── input5.txt
│   ├── output1.txt
│   ├── output2.txt
│   └── output5.txt

```

You are given two coding files: `Driver.java` and `HW3_Student_Solution.java`. `Driver.java` takes the input file, parses it and creates an instance of the class `HW3_Student_Solution` and prints the result to the command line. You only need to update the `HW3_Student_Solution.java` file. You may write your own helper methods and data structures in it.

The testcases folder has 3 input files and their corresponding output files for your reference. We will use these three input files (and seven others) in our autograding.

Method you need to write:

```

5.      /**
        * This method must be filled in by you. You may add other methods and subcl
        * but they must remain within the HW3_Student_Solution class.
        * @return Your resulting vector.
        */
        public int[] outputVector() {
            out_vector = in_vector;

            return out_vector;
        }
10.

```

The `HW3_Student_Solution` class has 2 instance variables.

- `in_vector` which is an array of ints, and stores the input vector given by the driver.
- `out_vector` which should be an array of ints, and will store your answer.

Compiling and executing from command line:

Assuming you're in the same directory level as `Driver.java`. Run `javac Driver.java` to compile.

To execute your code on `input1.txt`, run `java Driver testcases/input1.txt`. You can compare your output to `output1.txt`.

Submission

You only need to submit `HW3_Student_Solution.java` to Autolab.

Grading Guidelines

We will follow the usual grading guidelines for programming questions ([../policies/hw-policy.html#grading](#)).

Question 2 (Structured Upper Triangular matrix) [45 points]

The Problem

As we saw in the support page on matrix vector multiplication ([../support/matrix-vect/index.html](#)) (and in Question 1), if our matrix \mathbf{A} is guaranteed to have some structure, then we can *potentially* solve the matrix-vector multiplication faster than the more general case of arbitrary \mathbf{A} above. (Note that the vector \mathbf{x} remains arbitrary though.) We explore this phenomenon more in this problem and the next.

Given a vector \mathbf{r} of length n we define an $n \times n$ matrix $U^{\mathbf{r}}$ as follows. For any $0 \leq i, j < n$:

$$U^{\mathbf{r}}[i][j] = \begin{cases} 0 & \text{if } j < i \\ r[i] & \text{otherwise.} \end{cases}$$

For example, if $\mathbf{r} = (6, 17, 9)$, we have

$$U^{(6,17,9)} = \begin{pmatrix} 6 & 6 & 6 \\ 0 & 17 & 17 \\ 0 & 0 & 9 \end{pmatrix}.$$

Finally, we come to the actual question:

1. Design an $O(n)$ algorithm, which given any two vectors \mathbf{r} and \mathbf{x} of length n correctly computes $U^{\mathbf{r}} \cdot \mathbf{x}$.

Hint

Try and see how you can compute $y[i + 1]$ given $y[i]$ (where $\mathbf{y} = U^{\mathbf{r}} \cdot \mathbf{x}$.) Alternatively, if it helps, you can assume that there exists an algorithm to solve Question 1 and you can use it as a black box.

2. Argue that your algorithm in part (a) above runs in time $\Omega(n)$. (Hence, conclude that your algorithm runs in time $\Theta(n)$.)

Submission

You need to submit **one PDF** file to Autolab. We recommend that you typeset your solution but we will accept scans of handwritten solution-- you have to make sure that the scan is legible. Also make sure that you preview your upload on Autolab to make sure it was uploaded correctly.

Grading Guidelines

We will follow the usual grading guidelines for non-programming questions ([../policies/hw-policy.html#grading](http://www.cse.buffalo.edu/~atri/cse331/fall16/policies/hw-policy.html#grading)). Here is a high level grading rubric specific to this problem:

1. Algorithm idea: 18 points.
2. Algorithm details: 17 points.
3. Big-Oh analysis: 5 points.
4. Big-Omega analysis: 5 points.
5. Note: No proof idea or proof details of correctness is needed for this question. However your algorithm idea should implicitly include the proof idea of correctness. For an example of such an algorithm idea, see the algorithm idea for the structured matrix vector multiplication (for outer product) in the support page on matrix vector multiplication ([../support/matrix-vect/index.html](http://www.cse.buffalo.edu/~atri/cse331/fall16/support/matrix-vect/index.html)).

! Note

If you do not have separated out algorithm idea and algorithm details, you will get a zero(0) irrespective of the technical correctness of your solution..

! Note

You will get a zero on the entire problem if you present an $\Omega(n^2)$ time algorithm.

Questions 3 (Approaching FFT) [15 points]

The Problem

For this problem we will assume that $n = m^2$ for some integer $m \geq 2$. Further for any integers $0 \leq i, j < n$ we will denote their m -ary representations by (i_0, i_1) and (j_0, j_1) (where $i_0, i_1, j_0, j_1 \in \{0, \dots, m-1\}^2$). In other words,

$$i = i_0 + i_1 \cdot m,$$

and

$$j = j_0 + j_1 \cdot m.$$

We are now ready to define our structured matrix \mathbf{A} . For every k, ℓ such that $0 \leq k + \ell \leq 1$, we are given $m \times m$ matrices $\mathbf{B}^{k,\ell}$. Then the matrix \mathbf{A} is defined as follows. For every $0 \leq i, j < n$,

$$\mathbf{A}[i][j] = \prod_{k,\ell: 0 \leq k+\ell \leq 1} \mathbf{B}^{k,\ell}[i_k][j_\ell].$$

Let us first illustrate the definition above with an example. Let $m = 2$. In this case say we are given the following three matrices

$$\mathbf{B}^{0,0} = \begin{pmatrix} 1 & 1 \\ 2 & 2 \end{pmatrix},$$

$$\mathbf{B}^{0,1} = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix},$$

and

$$\mathbf{B}^{1,0} = \begin{pmatrix} 5 & 6 \\ 7 & 8 \end{pmatrix}.$$

Then the final matrix is

$$\begin{aligned} \mathbf{A} &= \begin{pmatrix} 1 \cdot 1 \cdot 5 & 1 \cdot 1 \cdot 6 & 1 \cdot 2 \cdot 5 & 1 \cdot 2 \cdot 6 \\ 2 \cdot 3 \cdot 5 & 2 \cdot 3 \cdot 6 & 2 \cdot 4 \cdot 5 & 2 \cdot 4 \cdot 6 \\ 1 \cdot 1 \cdot 7 & 1 \cdot 1 \cdot 8 & 1 \cdot 2 \cdot 7 & 1 \cdot 2 \cdot 8 \\ 2 \cdot 3 \cdot 7 & 2 \cdot 3 \cdot 8 & 2 \cdot 4 \cdot 7 & 2 \cdot 4 \cdot 8 \end{pmatrix} \\ &= \begin{pmatrix} 5 & 6 & 10 & 12 \\ 30 & 36 & 40 & 48 \\ 7 & 8 & 14 & 16 \\ 42 & 48 & 56 & 64 \end{pmatrix}, \end{aligned}$$

where in the above, we index the rows of \mathbf{A} by $(i_0, i_1) \in \{0, 1\}^2$ and are ordered from "top" to "bottom" as $(0, 0), (0, 1), (1, 0), (1, 1)$ while the columns of \mathbf{A} are indexed by $(j_0, j_1) \in \{0, 1\}^2$ and are ordered from "left" to "right" as $(0, 0), (0, 1), (1, 0), (1, 1)$.

Let's unpack this a bit

Just to make sure everyone is on the same page, let me walk through one of the calculations above. Say we are interested in $\mathbf{A}[2][1]$, i.e. $i = 2$ and $j = 1$. Then we consider the 2-ary (or binary) representation of $i = (i_1, i_0) = (1, 0)$ and $j = (j_1, j_0) = (0, 1)$. In other words, we have $i_0 = 0, i_1 = 1, j_0 = 1$ and $j_1 = 0$. Now note that by definition we have

$$\mathbf{A}[i][j] = \mathbf{B}^{0,0}[i_0][j_0] \cdot \mathbf{B}^{0,1}[i_0][j_1] \cdot \mathbf{B}^{1,0}[i_1][j_0].$$

In particular, we have

$$\mathbf{A}[2][1] = \mathbf{B}^{0,0}[0][1] \cdot \mathbf{B}^{0,1}[0][0] \cdot \mathbf{B}^{1,0}[1][1].$$

Now, note that we have $\mathbf{B}^{0,0}[0][1] = 1$, $\mathbf{B}^{0,1}[0][0] = 1$ and $\mathbf{B}^{1,0}[1][1] = 8$ and thus, we have

$$\mathbf{A}[2][1] = 1 \cdot 1 \cdot 8 = 8,$$

as stated above.

It is legitimate to wonder why should one look at this somewhat complicated way of defining such a matrix. Turns out that this family of matrices contains perhaps the most widely used structured matrix: the so called discrete Fourier matrix. (See the section on *Discrete Fourier Transform* in support page on matrix vector multiplication (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/hws/hw3/index.html>) for more on these matrices including their definition.)

Interestingly, this is probably one of the very few cases in CSE 331 that looking at a definition abstractly *might* be more beneficial than trying out some examples. The latter is definitely true when trying to design the algorithm in the question below.¹

So now to the actual problem. Design an algorithm that given arbitrary matrices $\mathbf{B}^{k,\ell}$ for every k, ℓ such that $0 \leq k + \ell \leq 1$ and an arbitrary vector \mathbf{x} of length n , compute $\mathbf{A} \cdot \mathbf{x}$, where \mathbf{A} is as defined above. Your algorithm should run in $O(nm) = O(n^{3/2})$ time.

Hint

The *distributive law strategy* outlined in support page on matrix vector multiplication (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/hws/hw3/index.html>) does lead to the desired algorithm. However, note that the above question is bit more involved than the example use of the distributive law strategy outlined in support page on matrix vector multiplication (<http://www-student.cse.buffalo.edu/~atri/cse331/fall16/hws/hw3/index.html>). It *might* also be useful to think of both \mathbf{x} and \mathbf{y} as an $m \times m$ matrix instead of just a traditional one-dimensional array of length $n = m^2$. So e.g. let \mathbf{X} be the $m \times m$ matrix such that $\mathbf{x}[i] = \mathbf{X}[i_0][i_1]$.

Note

it turns out that the above idea can be generalized to the case when $n = 2^m$ and results in an $O(n \log n)$ time algorithm to compute the DFT. However, you only need to show the weaker statement above.

Submission

You need to submit **one PDF** file to Autolab. We recommend that you typeset your solution but we will accept scans of handwritten solution-- you have to make sure that the scan is legible. Also make sure that you preview your upload on Autolab to make sure it was uploaded correctly.

Grading Guidelines

We will follow the usual grading guidelines for non-programming questions ([../policies/hw-policy.html#grading](http://www.cse.buffalo.edu/~atri/cse331/fall16/policies/hw-policy.html#grading)). Here is a high level grading rubric specific to this problem:

1. Algorithm idea: 6 points.
2. Algorithm details: 6 points.
3. Proof idea: 1 points for a proof idea that your algorithm is correct (proof details are not required).
4. Runtime analysis: 2 points for showing that your algorithm runs in $O(m^3)$ time.
5. Note: You will get a zero if your algorithm's runtime is $\Omega(n^2) = \Omega(m^4)$.

! Note

If you do not have separated out proof/algorithm idea and algorithm details, you will get a zero(0) irrespective of the technical correctness of your solution..